



Research Trends Analysis on AI and ML of 2026

Introduction

The "**Research Trends Analysis on AI and ML of 2026**" is a comprehensive Jupyter Notebook designed to build an **end-to-end intelligence pipeline** for the automated extraction and analysis of academic literature from the **ArXiv repository**. The system processes raw data through several sophisticated layers—ranging from **batch data acquisition** and **NLP-based cleaning** to **semantic embedding** and **LLM-powered research agent capabilities**. By transforming unstructured XML metadata into an enriched, actionable dataset, the notebook provides a robust framework for navigating the rapidly evolving landscape of Artificial Intelligence and Machine Learning.

Objective

The primary objective of this notebook is to **systematize the discovery of emerging research trends** and **identify high-signal papers** within the massive influx of daily publications. Key goals include:

- **Automated Data Gathering:** Fetching large-scale datasets (e.g., 7,000+ records) from the ArXiv API while respecting rate limits.
- **Semantic Intelligence:** Using **Sentence Transformers** and **K-Means clustering** to map the "true meaning" of papers beyond simple keywords, enabling precise semantic search and similar-paper retrieval.
- **Topic Modeling:** Applying **Non-negative Matrix Factorization (NMF)** to categorize research into core themes such as "Agents & Tool Use," "Memory & Attention Systems," and "Reinforcement Learning".
- **Automated Research Scoring:** Developing a **Research Agent System** that computes **novelty and impact scores** for each paper by measuring semantic distance from the corpus centroid and detecting high-value signals like "state of the art" or "open source".
- **Deep Insight Extraction:** Utilizing a **local LLM (Flan-T5)** to automatically extract core contributions, identified problems, and key methods from top-ranking papers to facilitate rapid literature review.

```
In [2]: import requests
import pandas as pd
import time #Used here to add delays between requests, respecting API rate lin
import xml.etree.ElementTree as ET
```

```

#Parses XML responses from the ArXiv API. ArXiv returns Atom XML feeds, so this
from tqdm import tqdm
#Adds a progress bar to loops, making it easier to track progress when fetching

BASE_URL = "http://export.arxiv.org/api/query"

'''
Defines a function to fetch papers from ArXiv.

Parameters:

search_query: What to search for (default is "machine learning").

max_results: Total number of papers to fetch.

batch_size: How many papers to fetch per request (ArXiv limits batch sizes).

delay: Pause between requests to avoid overwhelming the API.
'''
def fetch_arxiv_data(search_query="all:machine learning",
                    max_results=7000,
                    batch_size=100,
                    delay=3):

    """
    Fetch data from ArXiv API in batches.
    """

    all_entries = [] #Initializes an empty list to store all paper entries.

    '''
    Loops from 0 up to max_results in steps of batch_size.

    tqdm wraps the loop to show a progress bar.

    Example: If max_results=1000 and batch_size=100, it will run 10 iterations
    '''
    for start in tqdm(range(0, max_results, batch_size)):
        params = {
            "search_query": search_query,
            "start": start,
            "max_results": batch_size,
            "sortBy": "submittedDate",
            "sortOrder": "descending"
        }

        response = requests.get(BASE_URL, params=params)

        if response.status_code != 200:
            print(f"Error at start={start}: {response.status_code}")
            continue

```

```

'''
Converts the XML response into a tree structure.

Defines the Atom namespace (needed because XML tags are namespaced).
'''
root = ET.fromstring(response.content)

namespace = {"atom": "http://www.w3.org/2005/Atom"}

for entry in root.findall("atom:entry", namespace):
    paper = {}

    paper["id"] = entry.find("atom:id", namespace).text
    paper["title"] = entry.find("atom:title", namespace).text.strip()
    paper["summary"] = entry.find("atom:summary", namespace).text.strip()

    paper["published"] = entry.find("atom:published", namespace).text
    paper["updated"] = entry.find("atom:updated", namespace).text

    # Authors
    authors = entry.findall("atom:author", namespace)
    paper["authors"] = ", ".join([
        author.find("atom:name", namespace).text for author in authors
    ])

    # Categories
    '''
    Extracts subject categories (e.g., cs.AI, cs.LG).

    Joins them into a comma-separated string.
    '''
    categories = entry.findall("atom:category", namespace)
    paper["categories"] = ", ".join([
        cat.attrib["term"] for cat in categories
    ])

    all_entries.append(paper)

    time.sleep(delay) # respect API rate limits

return pd.DataFrame(all_entries)

def save_to_csv(df, filename="arxiv_data.csv"):
    df.to_csv(filename, index=False)
    print(f"Saved {len(df)} records to {filename}")

if __name__ == "__main__":
    df = fetch_arxiv_data(
        search_query="cat:cs.AI OR cat:cs.LG OR cat:cs.CL",
        max_results=7000,
        batch_size=100
    )

```

```
)  
  
save_to_csv(df)
```

100%|██████████| 70/70 [09:53<00:00, 8.48s/it]

Saved 7000 records to arxiv_data.csv

Data Cleaning + NLP Pipeline

```
In [3]: import pandas as pd  
import re  
import nltk  
import spacy  
import yake  
from tqdm import tqdm  
  
# Download NLTK resources (run once)  
nltk.download('stopwords')  
from nltk.corpus import stopwords  
  
# Load spaCy model  
nlp = spacy.load("en_core_web_sm")  
  
STOPWORDS = set(stopwords.words('english'))  
  
# =====  
# 1. LOAD DATA  
# =====  
file_path = r"C:\Users\abc\Downloads\arxiv_data.csv"  
#df = pd.read_csv(file_path)  
df = pd.read_csv(file_path, encoding='utf-8', encoding_errors='ignore')  
  
# =====  
# 2. CLEANING FUNCTIONS  
# =====  
  
def clean_text(text):  
    if pd.isna(text):  
        return ""  
  
    text = text.lower()  
  
    # Remove URLs  
    text = re.sub(r"http\S+", "", text)  
  
    # Remove special characters / numbers  
    text = re.sub(r"[^a-zA-Z\s]", " ", text)  
  
    # Remove extra spaces  
    text = re.sub(r"\s+", " ", text).strip()  
  
    return text
```

```

# Apply cleaning
df["clean_summary"] = df["summary"].apply(clean_text)
df["clean_title"] = df["title"].apply(clean_text)

# =====
# 3. TOKENIZATION + LEMMATIZATION
# =====

'''
Uses spaCy to tokenize text.

Keeps only alphabetic tokens.

Removes stopwords.

Filters out very short words (length ≤ 2).

Converts words to their lemma (base form, e.g., "running" → "run").
'''
def tokenize(text):
    doc = nlp(text)
    tokens = [
        token.lemma_
        for token in doc
        if token.is_alpha
        and token.text not in STOPWORDS
        and len(token.text) > 2
    ]
    return tokens

tqdm.pandas()
df["tokens"] = df["clean_summary"].progress_apply(tokenize)

# =====
# 4. KEYWORD EXTRACTION (YAKE)
# =====

'''
Configures YAKE to extract keywords.

n=2 → allows up to bigrams (two-word phrases).

top=5 → extracts top 5 keywords per document.
'''
kw_extractor = yake.KeywordExtractor(
    lan="en",
    n=2,           # max n-gram size
    top=5          # top keywords
)

```

```

def extract_keywords(text):
    if not text:
        return []
    keywords = kw_extractor.extract_keywords(text)
    return [kw[0] for kw in keywords]

df["keywords"] = df["clean_summary"].progress_apply(extract_keywords)

# =====
# 5. CATEGORY CLEANING
# =====

def extract_primary_category(cat):
    if pd.isna(cat):
        return None
    return cat.split(",")[0].strip()

df["primary_category"] = df["categories"].apply(extract_primary_category)

# =====
# 6. DATE FEATURES
# =====

df["published"] = pd.to_datetime(df["published"])
df["year"] = df["published"].dt.year
df["month"] = df["published"].dt.month

# =====
# 7. SAVE CLEAN DATA
# =====

output_path = r"C:\Users\abc\Downloads\arxiv_cleaned.csv"
df.to_csv(output_path, index=False)

print("✅ Pipeline complete. Cleaned data saved.")

```

```

[nltk_data] Downloading package stopwords to
[nltk_data]      C:\Users\abc\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
100%|██████████| 7000/7000 [09:40<00:00, 12.05it/s]
100%|██████████| 7000/7000 [02:07<00:00, 54.81it/s]
✅ Pipeline complete. Cleaned data saved.

```

Semantic + Trend Intelligence Layer

STEP 1 — Load & Validate (Production-safe)

```

In [4]: import pandas as pd
import ast
# Provides literal_eval, which safely evaluates strings that look like Python

```

```

# #(like lists or dictionaries) without executing arbitrary code.

file_path = r"C:\Users\abc\Downloads\arxiv_cleaned.csv"

df = pd.read_csv(file_path, encoding="utf-8")

# Convert stringified lists safely
'''
Defines a helper function to convert stringified lists back into actual Python
Uses ast.literal_eval (safe, unlike eval) to parse strings like "['AI', 'learn
If parsing fails (e.g., malformed string), it returns an empty list instead of
'''
def safe_eval(x):
    try:
        return ast.literal_eval(x)
    except:
        return []

#This is crucial because when saving to CSV, lists are stored as strings.
# Without conversion, you couldn't easily manipulate them later.
df["tokens"] = df["tokens"].apply(safe_eval)
df["keywords"] = df["keywords"].apply(safe_eval)

# Basic sanity checks
assert "clean_summary" in df.columns
assert df["clean_summary"].notna().sum() > 0

print("Dataset shape:", df.shape)

```

Dataset shape: (7000, 14)

STEP 2 — Build Document Representation (TF-IDF)

```

In [5]: '''TF-IDF = Term Frequency–Inverse Document Frequency,
a common technique in NLP to measure how important a word is in a document
relative to the whole corpus.
'''

from sklearn.feature_extraction.text import TfidfVectorizer

'''
Creates a TF-IDF vectorizer with specific settings:

max_features=5000 → Keep only the 5,000 most important terms. Prevents the mat
ngram_range=(1,2) → Includes both single words (unigrams) and two-word phrases
'''

```

```

stop_words="english" → Removes common English stopwords (like "the", "and").

min_df=5 → Ignores terms that appear in fewer than 5 documents. Helps remove r

max_df=0.8 → Ignores terms that appear in more than 80% of documents. Prevents
...

tfidf = TfidfVectorizer(
    max_features=5000,
    ngram_range=(1,2),
    stop_words="english",
    min_df=5,
    max_df=0.8
)

X_tfidf = tfidf.fit_transform(df["clean_summary"])

print("TF-IDF shape:", X_tfidf.shape)

...
fit_transform does two things:

Fit → Learns the vocabulary and IDF weights from the corpus (clean_summary).

Transform → Converts each summary into a sparse TF-IDF vector.

X_tfidf is a sparse matrix where:

Rows = documents (summaries).

Columns = features (words/phrases).

Values = TF-IDF scores.
...

```

TF-IDF shape: (7000, 5000)

Out[5]: '\nfit_transform does two things:\n\nFit → Learns the vocabulary and IDF weights from the corpus (clean_summary).\n\nTransform → Converts each summary into a sparse TF-IDF vector.\n\nX_tfidf is a sparse matrix where:\n\nRows = documents (summaries).\n\nColumns = features (words/phrases).\n\nValues = TF-IDF scores.\n'

```

In [6]: import matplotlib.pyplot as plt    # <-- add this line
import seaborn as sns
from sklearn.manifold import TSNE


tsne = TSNE(n_components=2, random_state=42, perplexity=30)
X_tsne = tsne.fit_transform(X_tfidf.toarray())

plt.figure(figsize=(8,6))
plt.scatter(X_tsne[:,0], X_tsne[:,1], alpha=0.5)
plt.title("t-SNE visualization of TF-IDF features")
plt.show()


```



```
sns.scatterplot(  
    x=X_tsne[:,0], y=X_tsne[:,1],  
    hue=df["primary_category"],  
    palette="tab10",  
    alpha=0.7  
)  
plt.title("t-SNE visualization colored by category")  
plt.show()
```



t-SNE visualization colored by category



STEP 3 — Topic Modeling (Core Research Insight)

```
In [7]: from sklearn.decomposition import NMF
#Brings in Non-negative Matrix Factorization (NMF) from scikit-learn.
#NMF is a dimensionality reduction technique often used for topic modeling in

n_topics = 10
#Creates an NMF model with n_components=10 (i.e., 10 topics).

#random_state=42 ensures reproducibility (same results every run).

nmf = NMF(n_components=n_topics, random_state=42)
W = nmf.fit_transform(X_tfidf)
H = nmf.components_

'''fit_transform does two things:

Learns the topic structure from the TF-IDF matrix (X_tfidf).

Produces the document-topic matrix W.

W has shape (n_documents, n_topics).

Each row = a document.

Each column = how strongly that document belongs to a topic.
'''



feature_names = tfidf.get_feature_names_out()
```



```
In [8]: import matplotlib.pyplot as plt



n_top_words = 10



for topic_idx, topic in enumerate(H):
    top_indices = topic.argsort()[-n_top_words:]
    top_words = [feature_names[i] for i in top_indices]
    top_weights = topic[top_indices]



    plt.figure(figsize=(8,4))
    plt.barh(top_words, top_weights)
    plt.title(f"Topic {topic_idx}")
    plt.xlabel("Weight")
    plt.show()
```











Extract Topics (VERY IMPORTANT STEP)

In [9]: '''Defines a function called display_topics.

Inputs:

H → the topic-term matrix from NMF (nmf.components_). Each row = a topic, each
 feature_names → the vocabulary list from TF-IDF (tfidf.get_feature_names_out())
 n_top_words → how many top words to show per topic (default = 10).


```

...

def display_topics(H, feature_names, n_top_words=10):
    topics = {}
    for topic_idx, topic in enumerate(H):
        top_words = [feature_names[i] for i in topic.argsort()[: -n_top_words -
        topics[topic_idx] = top_words
        print(f"\nTopic {topic_idx}:")
        print(", ".join(top_words))
    return topics

topics = display_topics(H, feature_names)

```

Topic 0:
llms, models, language, llm, language models, large, large language, model, models llms, evaluation

Topic 1:
neural, gradient, space, algorithm, linear, problem, quantum, optimization, error, order

Topic 2:
agent, agents, multi agent, multi, execution, tool, llm, agentic, skill, task

Topic 3:
rl, policy, reward, reinforcement, learning, reinforcement learning, training, optimization, rewards, learning rl

Topic 4:
retrieval, rag, graph, knowledge, evidence, retrieval augmented, augmented, generation, query, augmented generation

Topic 5:
memory, attention, token, inference, kv, tokens, compression, context, layer, long

Topic 6:
reasoning, cot, thought, chain, models, step, traces, chain thought, answer, accuracy

Topic 7:
visual, multimodal, image, vision, video, text, semantic, vlms, vision language, spatial

Topic 8:
data, learning, detection, real, time, model, datasets, based, performance, dataset

Topic 9:
ai, human, systems, research, ai systems, governance, intelligence, design, artificial, users

STEP 4 — Assign Topic to Each Paper

```
In [10]: df["topic"] = W.argmax(axis=1)
```

STEP 5 — Trend Analysis


```
In [11]: import matplotlib.pyplot as plt

# Example: manually assign labels based on top words
topic_labels = {
    0: "ML & Prediction Systems",
    1: "Agents & Tool Use",
    2: "Multimodal / Vision",
    3: "Reinforcement Learning",
    4: "AI Systems & Evaluation",
    5: "Multilingual & Speech",
    6: "Reasoning & CoT",
    7: "Knowledge & Continual Learning",
    8: "Neural Theory / Foundations",
    9: "Memory & Attention Systems"
}

# Replace numeric topic IDs with labels
df["topic_label"] = df["topic"].map(topic_labels)

# Count by label
topic_counts = df["topic_label"].value_counts()

# Plot with labels
plt.figure(figsize=(10,5))
topic_counts.plot(kind="bar")
plt.title("Distribution of Research Topics")
plt.xlabel("Topic")
plt.ylabel("Number of Papers")
plt.xticks(rotation=45, ha="right")
plt.show()
```



Quick Visualization


```
In [12]: import matplotlib.pyplot as plt

# Example: assign descriptive names to topics
topic_labels = {
    0: "ML & Prediction Systems",
    1: "Agents & Tool Use",
    2: "Multimodal / Vision",
    3: "Reinforcement Learning",
    4: "AI Systems & Evaluation",
    5: "Multilingual & Speech",
    6: "Reasoning & CoT",
    7: "Knowledge & Continual Learning",
    8: "Neural Theory / Foundations",
    9: "Memory & Attention Systems"
}

# Add a new column with human-readable labels
df["topic_label"] = df["topic"].map(topic_labels)

# Count by label instead of numeric index
topic_counts = df["topic_label"].value_counts()
```

```
# Plot with descriptive labels
plt.figure(figsize=(10,5))
topic_counts.plot(kind="bar")
plt.title("Distribution of Research Topics")
plt.xlabel("Topic")
plt.ylabel("Number of Papers")
plt.xticks(rotation=45, ha="right")
plt.show()
```



STEP 6 — Save Enriched Dataset

```
In [13]: output_path = r"C:\Users\abc\Downloads\arxiv_with_topics.csv"
df.to_csv(output_path, index=False)
```

Temporal Dynamics + Emerging Trend Detection

```
In [14]: # =====
# 0. IMPORTS
# =====
import pandas as pd
import ast
from collections import Counter
import matplotlib.pyplot as plt

# =====
# 1. LOAD DATA (ROBUST)
# =====
file_path = r"C:\Users\abc\Downloads\arxiv_with_topics.csv"

df = pd.read_csv(file_path, encoding="utf-8")

# =====
# 2. SAFE PARSING (CRITICAL)
# =====
def safe_eval(x):
    try:
        return ast.literal_eval(x) if isinstance(x, str) else []
    except:
        return []

df["keywords"] = df["keywords"].apply(safe_eval)
df["tokens"] = df["tokens"].apply(safe_eval)

# =====
# 3. DATA VALIDATION
# =====
assert "topic" in df.columns, "Topic column missing"
assert df.shape[0] > 0, "Empty dataset"

print("✅ Dataset loaded:", df.shape)

# =====
# 4. HUMAN-READABLE TOPIC LABELS
# =====
topic_labels = {
    0: "ML & Prediction Systems",
    1: "Agents & Tool Use",
    2: "Multimodal / Vision",
    3: "Reinforcement Learning",
    4: "AI Systems & Evaluation",
```

```

5: "Multilingual & Speech",
6: "Reasoning & CoT",
7: "Knowledge & Continual Learning",
8: "Neural Theory / Foundations",
9: "Memory & Attention Systems"
}

df["topic_label"] = df["topic"].map(topic_labels)

# =====
# 5. TOPIC DISTRIBUTION
# =====
topic_dist = df["topic_label"].value_counts().sort_values(ascending=False)

print("\n📊 Topic Distribution:")
print(topic_dist)

# =====
# 6. KEYWORD SIGNAL EXTRACTION
# =====
def extract_top_keywords(df, topic_id, n=10):
    subset = df[df["topic"] == topic_id]
    keywords = []

    for row in subset["keywords"]:
        if isinstance(row, list):
            keywords.extend(row)

    return Counter(keywords).most_common(n)

print("\n🧠 Topic Keyword Signals:")
topic_keyword_map = {}

for t in sorted(df["topic"].unique()):
    top_kws = extract_top_keywords(df, t)
    topic_keyword_map[t] = top_kws

    print(f"\n{topic_labels[t]}")
    print(top_kws)

# =====
# 7. TOP PAPERS PER TOPIC
# =====
def get_top_papers(df, topic_id, n=5):
    subset = df[df["topic"] == topic_id]
    subset = subset.sort_values("published", ascending=False)

    return subset[["title", "authors"]].head(n)

print("\n📄 Representative Papers per Topic:")

for t in sorted(df["topic"].unique()):
    print(f"\n=== {topic_labels[t]} ===")

```

```

    print(get_top_papers(df, t).to_string(index=False))

# =====
# 8. GLOBAL HIGH-SIGNAL PAPERS
# =====
df["keyword_count"] = df["keywords"].apply(lambda x: len(x) if isinstance(x, list) else 0)

top_papers = df.sort_values("keyword_count", ascending=False).head(10)

print("\n🔥 High Signal Papers:")
print(top_papers[["title", "topic_label", "keyword_count"]].to_string(index=False))

# =====
# 9. CLEAN VISUALIZATION
# =====
plt.figure(figsize=(12,6))

topic_dist.plot(kind="bar")

plt.title("Research Topic Distribution (arXiv Snapshot)", fontsize=14)
plt.xlabel("Topic", fontsize=12)
plt.ylabel("Number of Papers", fontsize=12)

plt.xticks(rotation=45, ha="right")
plt.grid(axis="y", linestyle="--", alpha=0.7)

plt.tight_layout()
plt.show()

# =====
# 10. SAVE FINAL DATASET
# =====
output_path = r"C:\Users\abc\Downloads\arxiv_intelligence.csv"
df.to_csv(output_path, index=False)

print(f"\n✅ Saved enriched dataset to:\n{output_path}")

```

✅ Dataset loaded: (7000, 16)

📊 Topic Distribution:

topic_label	
Agents & Tool Use	1283
Neural Theory / Foundations	1255
ML & Prediction Systems	785
Multimodal / Vision	708
Multilingual & Speech	643
Knowledge & Continual Learning	638
Reinforcement Learning	478
Memory & Attention Systems	439
Reasoning & CoT	396
AI Systems & Evaluation	375

Name: count, dtype: int64

🧠 Topic Keyword Signals:

ML & Prediction Systems

[('large language', 164), ('language models', 63), ('fine tuning', 51), ('natural language', 19), ('real world', 16), ('cross lingual', 16), ('fine tuned', 14), ('llm based', 13), ('open source', 11), ('low resource', 11)]

Agents & Tool Use

[('neural networks', 52), ('neural network', 35), ('high dimensional', 25), ('physics informed', 22), ('gradient descent', 17), ('flow matching', 16), ('machine learning', 15), ('large language', 11), ('sample complexity', 11), ('deep learning', 11)]

Multimodal / Vision

[('multi agent', 92), ('large language', 63), ('real world', 32), ('llm agents', 25), ('llm based', 23), ('language models', 20), ('open source', 18), ('long horizon', 15), ('language model', 14), ('multi turn', 14)]

Reinforcement Learning

[('reinforcement learning', 98), ('large language', 29), ('fine tuning', 27), ('post training', 19), ('language models', 14), ('policy optimization', 13), ('group relative', 9), ('test time', 9), ('multi objective', 9), ('open source', 9)]

AI Systems & Evaluation

[('large language', 41), ('language models', 26), ('retrieval augmented', 24), ('augmented generation', 23), ('knowledge graph', 18), ('language model', 11), ('knowledge graphs', 11), ('question answering', 9), ('llm based', 9), ('external knowledge', 9)]

Multilingual & Speech

[('large language', 67), ('language models', 35), ('long context', 20), ('low rank', 20), ('fine tuning', 19), ('speculative decoding', 14), ('long term', 12), ('inference time', 11), ('language model', 10), ('long horizon', 9)]

Reasoning & CoT

[('large language', 60), ('language models', 27), ('reasoning traces', 13), ('test time', 13), ('reasoning models', 11), ('reinforcement learning', 10), ('tim

e scaling', 8), ('latent reasoning', 7), ('vision language', 7), ('mathematical reasoning', 7)]

Knowledge & Continual Learning

[('vision language', 65), ('cross modal', 39), ('language models', 36), ('fine grained', 31), ('fine tuning', 19), ('large language', 19), ('real world', 17), ('multi modal', 13), ('diffusion models', 12), ('multimodal large', 11)]

Neural Theory / Foundations

[('machine learning', 53), ('time series', 52), ('deep learning', 47), ('real world', 43), ('real time', 31), ('anomaly detection', 26), ('foundation models', 24), ('large scale', 23), ('federated learning', 20), ('data driven', 19)]

Memory & Attention Systems

[('artificial intelligence', 27), ('large language', 19), ('decision making', 12), ('language models', 12), ('mental health', 8), ('natural language', 6), ('real world', 6), ('large scale', 6), ('general purpose', 5), ('scientific discovery', 5)]

Representative Papers per Topic:

=== ML & Prediction Systems ===

title

authors

Representational Harms in LLM-Generated Narratives Against Global Majority Nationalities
Ilana Nguyen, Harini Suresh, Thema Monroe-White, Evan Shieh

Neural Recovery of Historical Lexical Structure in Bantu Languages from Modern Data
Hillary Mutisya, John Mugane

Dharma, Data and Deception: An LLM-Powered Rhetorical Analysis of Cow-Urine Health Claims on YouTube
Sheza Munir, Ratna Kandala, Anamta Khan, Deepti, Joyojeet Pal

From Natural Language to Verified Code: Toward AI Assisted Problem-to-Code Generation with Dafny-Based Formal Verification
Md Erfan, Md Kamal Hossain Chowdhury, Ahmed Ryan, Md Rayhanur Rahman

Controllable Spoken Dialogue Generation: An LLM-Driven Grading System for K-12 Non-Native English Learners
Haidong Yuan, Haokun Zhao, Wanshi Xu, Songjun Cao, Qingyu Zhou, Long Ma, Hongjie Fan

=== Agents & Tool Use ===

title

authors

Relaxation-Informed Training of Neural Network Surrogate Models
Calvin Tsay

An Undecidability Proof for the Plan Existence Problem
Antonis Achilleos

Time-Localized Parametric Decomposition of Respiratory Airflow for Sub-Breath Analysis
Victoria Ribeiro Rodrigues, Paul W. Davenport, Nicholas J. Napoli

Operational Feature Fingerprints of Graph Datasets

ts via a White-Box Signal-Subspace Probe Yuchen Xiong, Swe
e Keong Yeap, Zhen Hong Ban
Iterative Model-Learning Scheme via Gaussian Processes for Nonlinear Model Pred
ictive Control of (Semi-)Batch Processes Tai Xuan Tan, Al
exander Mitsos, Eike Cramer

=== Multimodal / Vision ===

title

authors

Agentic World Modeling: Foundations, Capabilities,
Laws, and Beyond Meng Chu, Xuan Billy Zhang, Kevin Qinghong Lin, Lingdong Kong,
Jize Zhang, Teng Tu, Weijian Ma, Ziqi Huang, Senqiao Yang, Wei Huang, Yeying Ji
n, Zhefan Rao, Jinhui Ye, Xinyu Lin, Xichen Zhang, Qisheng Hu, Shuai Yang, Leya
ng Shen, Wei Chow, Yifei Dong, Fengyi Wu, Quanyu Long, Bin Xia, Shaozuo Yu, Min
gkang Zhu, Wenhui Zhang, Jiehui Huang, Haokun Gui, Haoxuan Che, Long Chen, Qifen
g Chen, Wenxuan Zhang, Wenya Wang, Xiaojuan Qi, Yang Deng, Yanwei Li, Mike Zhen
g Shou, Zhi-Qi Cheng, See-Kiong Ng, Ziwei Liu, Philip Torr, Jiaya Jia

QuantClaw: Precision Where It Mat

ters for OpenClaw

Manyi Zhang, Ji-Fu Li, Zhongao Sun, Xiaohao Liu, Zhenhua Dong, Xianzhi Yu, Haol
i Bai, Xiaobo Xia

Superminds Test: Actively Evaluating Collective Intelligence of Agent Society v
ia Probing Agents

Xirui Li, Ming Li, Yunze Xiao, Ryan Wong, Dianqi Li, Timothy Baldwin, Tianyi Zh
ou

From Skills to Talent: Organising Heterogeneous Agents as a R
eal-World Company

Zhengxu Yu, Yu Fu, Zhiyuan He, Yuxuan Huang, Lee Ka Yiu, Meng Fang, Weilin Luo,
Jun Wang

AgentSearchBench: A Benchmark for AI Agent S

earch in the Wild

Bin Wu, Arastun Mammadli, Xiaoyu Zhang, Emine Yilmaz

=== Reinforcement Learning ===

title

authors

CRAFT: Clustered Regression for Adaptive Filterin
g of Training data

Parthasarathi Panda, Asheswari Swain, Subhrakanta Panda

Learning Evidence Highlighti
ng for Frozen LLMs Shaoang Li, Yanhang Shi, Yufei Li, Mingfu Liang, Xiaohan We
i, Yunchen Pu, Fei Tian, Chonglin Sun, Frank Shyu, Luke Simon, Sandeep Pandey,
Xi Liu, Jian Li

SOLAR-RL: Semi-Online Long-horizon Assignment Rein
forcement Learning Jichao Wang, Liuyang Bian, Yufeng Zhou, Han Xiao, Yue
Pan, Guozhi Wang, Hao Wang, Zhaoxiong Wang, Yafei Wen, Xiaoxin Chen, Shuai Ren,
Lingfang Zeng

QDTraj: Exploration of Diverse Trajectory Primitives for Articulated Objects Ro
botic Manipulation

Mathilde Kappel, Mahdi Khoramshahi, Louis Annabi, Faiz Ben Amar, Stéphane Donci
eux

Preserve Support, Not Correspondence: Dynamic Routing for Offline Rein

forcement Learning
Zhancun Mu, Guangyu Zhao, Yiwu Zhong, Chi Zhang

=== AI Systems & Evaluation ===

title
authors
Aligning Dense Retrievers with LLM Utility via Distillation
Aligning Dense Retrievers with LLM Utility via Distillation Rajinder Sandhu, Di Mu, Cheng Chang, Md Shahriar Tasjid, Himanshu Rai, Maksims Volkovs, Ga Wu
BERAG: Bayesian Ensemble Retrieval-Augmented Generation for Knowledge-based Visual Question Answering
Jinghong Chen, Jingbiao Mei, Guangyu Yang, Bill Byrne
Can QPP Choose the Right Query Variant? Evaluating Query Variant Selection for RAG Pipelines
Negar Arabzadeh, Andrew Drozdov, Michael Bendersky, Matei Zaharia
How Hard is it to Decide if a Fact is Relevant to a Query?
Meghyn Bienvenu, Diego Figueira, Pierre Lafourcade
Distance-Misaligned Training in Graph Transformers and Adaptive Graph-Aware Control
Qinhan Hou, Jing Tang

=== Multilingual & Speech ===

title
authors
Spend Less, Fit Better: Budget-Efficient Scaling Law Fitting via Active Experiment Selection
Sijie Li, Shanda Li, Haowei Lin, Weiwei Sun, Ameet Talwalkar, Yiming Yang
How Do AI Agents Spend Your Money? Analyzing and Predicting Token Consumption in Agentic Coding Tasks
Longju Bai, Zhemin Huang, Xingyao Wang, Jiao Sun, Rada Mihalcea, Erik Brynjolfsson, Alex Pentland, Jiaxin Pei
Adaptive Head Budgeting for Efficient Multi-Head Attention
Bilal Faye, Abdoulaye Mbaye, Hanane Azzag, Mustapha Lebbah
SpikingBrain2.0: Brain-Inspired Foundation Models for Efficient Long-Context and Cross-Platform Inference
Yuqi Pan, Jinghao Zhuang, Yupeng Feng, Fangzhi Zhong, Siyu Ding, Xuerui Qiu, Shaowei Gu, Bohan Sun, Zhiyong Qin, Yibo Zhong, Lingtao Ouyang, Kun Yang, Zehao Liu, Yuhong Chou, Shurong Wang, Anjie Hu, Han Xu, Bo Xu, Guoqi Li
RouteLMT: Learned Sample Routing for Hybrid LLM Translation Deployment
Yingfeng Luo, Hongyu Liu, Dingyang Lin, Kaiyan Chang, Chenglong Wang, Bei Li, Qian Du, Tong Xiao, Jingbo Zhu

=== Reasoning & CoT ===

title
authors
Thinking Without Words: Efficient Latent Reasoning with Abstract Chain-of-Thought
Keshav Ramji, Tahira Naseem, Ramón Fernandez Astudillo

Rethinking Math Reasoning Evaluation: A Robust LLM-as-a-Judge Framework Beyond Symbolic Rigidity Erez Yosef, Oron Anschel, Shunit Haviv Hakimi, Asaf Gendler, Adam Botach, Nimrod Berman, Igor Kviatkovsky

SSG: Logit-Balanced Vocabulary Partitioning for LLM Watermarking

Chenxi Gu, Xiaoning Du, John Grundy

Large Language Models

Decide Early and Explain Later Ayan Datta, Zhixue Zhao, Bhuvanesh Verma, Radhika Mamidi, Mounika Marreddy, Alexander Mehler
Towards Safe Mobility: A Unified Transportation Foundation Model enabled by Open-Ended Vision-Language Dataset Wenhui Huang, Songyan Zhang, Collister Chua, Yang Liang, Zhiqi Mao, Heng Yang, Chen Lv

=== Knowledge & Continual Learning ===

title

authors

Cross-Stage Coherence in Hierarchical Driving VQA: Explicit Baselines and Learned Gated Context Projectors
Gautam Kumar Jain, Carsten Markgraf, Julian Stähler

CGC: Compositional

Grounded Contrast for Fine-Grained Multi-Image Understanding
Lihao Zheng, Zhenwei Shao, Yu Zhou, Yan Yang, Xintian Shen, Jiawei Chen, Hao Ma, Tao Wei

Contrastive Semantic

Projection: Faithful Neuron Labeling with Contrastive Examples
Oussama Bouanani, Jim Berend, Wojciech Samek, Sebastian Lapuschkin, Maximilian Dreyer

Selective

Contrastive Learning For Gloss Free Sign Language Translation
Changhao Lai, Rui Zhao, Xuewen Zhong, Jinsong Su, Yidong Chen
ChangeQuery: Advancing Remote Sensing Change Analysis for Natural and Human-Induced Disasters from Visual Detection to Semantic Understanding Dongwei Sun, Jing Yao, Kan Wei, Xiangyong Cao, Chen Wu, Zhenghui Zhao, Pedram Ghamisi, Jun Zhou, Jón Atli Benediktsson

=== Neural Theory / Foundations ===

title

author

Zero-Shot Morphological Discovery in Low-Resource Bantu Languages via Cross-Lingual Transfer and Unsupervised Clustering
Hillary Mutisya, John Mugane

Quality-Driven

Selective Mutation for Deep Learning Zaheed Ahmed, Emmanuel Charleson Dapaah, Philip Makedonski, Jens Grabowski

Adversarial Malware Generation in Linux ELF Binaries via Semantic-Preserving Transformations
Lukáš Hrdonka, Martin Jureček

CLVAE: A Variational Autoencoder for

Long-Term Customer Revenue Forecasting Jeffrey Näf, Riana Valera Mbelson, Markus Meierer

Detecting Concept Drift in Evolving Malware Families Using Rule-Based Classifier Representations To

máš Kalný, Martin Jureček, Mark Stamp

=== Memory & Attention Systems ===

title

authors

How Supply Chain Dependencies Complicate Bias Measurement and Accountability Attribution in AI Hiring Applications

Gauri Sharma, Maryam Molamohammadi

Rethinking XAI Evaluation: A Human-Centered Audit of Shapley Benchmarks in High-Stakes Settings Inês Oliveira e Silva, Sérgio Jesus, Iker Perez, Rita P. Ribeiro, Carlos Soares, Hugo Ferreira, Pedro Bizarro

Measuring and Mitigating Persona Distortions from AI Writing Assistance Paul Röttger, Kobi Hackenburg, Hannah Rose Kirk, Christopher Summerfield

On the Hybrid Nature of ABPMS Process Frames and its Implications on Automated Process Discovery Anti Al

man, Izack Cohen, Avigdor Gal, Fabrizio Maria Maggi, Marco Montali

A Co-Evolutionary Theory of Human-AI Coexistence: Mutualism, Governance, and Dynamics in Complex Societies

Somyajit Chakraborty

🔥 High Signal Papers:

title

topic_label keyword_count

Spend Less, Fit Better: Budget-Efficient Scaling Law Fitting via Active Experiment Selection Multilingual & Speech 5

How Do AI Agents Spend Your Money? Analyzing and Predicting Token Consumption in Agentic Coding Tasks Multilingual & Speech 5

Representational Harms in LLM-Generated Narratives Against Global Majority Nationalities ML & Prediction Systems 5

Agentic World Modeling: Foundations, Capabilities, Laws, and Beyond Multimodal / Vision 5

Relaxation-Informed Training of Neural Network Surrogate Models Agents & Tool Use 5


An Undecidability Proof for the Plan Existence Problem Agents & Tool Use 5

Neural Recovery of Historical Lexical Structure in Bantu Languages from Modern Data ML & Prediction Systems 5

Zero-Shot Morphological Discovery in Low-Resource Bantu Languages via Cross-Lingual Transfer and Unsupervised Clustering Neural Theory / Foundations 5

Aligning Dense Retrievers with LLM Utility via Distillation Aligning Dense Retrievers with LLM Utility via Distillation AI Systems & Evaluation 5

Thinking Without Words: Efficient Latent Reasoning with Abstract Chain-of-Thought Reasoning & CoT



✓ Saved enriched dataset to:
C:\Users\abc\Downloads\arxiv_intelligence.csv

SEMANTIC INTELLIGENCE

We're going to build:

- ✓ 1. Embeddings (true meaning, not keywords)
- ✓ 2. Semantic search (query → relevant papers)
- ✓ 3. Similar paper retrieval (paper → neighbors)
- ✓ 4. Trend signal layer (emerging ideas detection)

1. LOAD + PREP

```
In [15]: import pandas as pd
import ast

file_path = r"C:\Users\abc\Downloads\arxiv_intelligence.csv"
df = pd.read_csv(file_path)

def safe_eval(x):
    try:
        return ast.literal_eval(x) if isinstance(x, str) else []
    except:
        return []

df["tokens"] = df["tokens"].apply(safe_eval)

# Combine meaningful text (IMPORTANT)
```

```
df["text"] = df["clean_title"] + " " + df["clean_summary"]  
'''
```

Creates a new column text by concatenating the cleaned title and summary.

This gives you a single unified text field per paper, which is very useful for Embedding models (e.g., sentence transformers).

Topic modeling.

Classification tasks.

Search/retrieval pipelines.

```
'''
```

```
print("✅ Loaded:", df.shape)
```

✅ Loaded: (7000, 18)

2. EMBEDDINGS (CORE UPGRADE)

The purpose of this step is to build the "**Semantic Intelligence Layer**" of the pipeline by converting textual research data into **mathematical representations known as embeddings**.

While previous steps in the notebook used keyword-based techniques like TF-IDF, this step uses a **Sentence Transformer model** (`all-MiniLM-L6-v2`) to capture the "**true meaning**" and context of the papers rather than just matching specific words.

Specific Components of the Step

- **Model Selection:** It initializes the `all-MiniLM-L6-v2` model, which is described as a **fast and strong baseline** for generating these dense vectors.
- **Data Input:** It converts the combined "text" column (which contains both the cleaned title and summary) into a list of strings to be processed.
- **Vector Generation:** The `model.encode()` function transforms each paper's text into a **dense vector (embedding)**, processing them in batches of 32 for efficiency.

Why These Embeddings Are Critical to the Notebook

According to the sources, generating these embeddings is a "core upgrade" that

enables several advanced features:

- **Semantic Search:** It allows you to search for papers using natural language queries by converting the query into an embedding and finding papers with the highest **cosine similarity**.
- **Similar Paper Retrieval:** By comparing these vectors, the system can identify "neighbors"—papers that are semantically similar to a specific document.
- **Novelty Scoring:** The notebook uses these embeddings to calculate a **Novelty Score** by measuring the "semantic distance" between an individual paper and the average (centroid) of the entire dataset.
- **Advanced Clustering:** These embeddings serve as the input for **K-Means clustering**, which groups papers into distinct research themes based on their semantic content.

```
In [16]: from sentence_transformers import SentenceTransformer

model = SentenceTransformer("all-MiniLM-L6-v2") # fast + strong baseline

# Generate embeddings
embeddings = model.encode(
    df["text"].tolist(),
    show_progress_bar=True,
    batch_size=32
)
...
df["text"].tolist() → converts your combined title+summary column into a list
model.encode() → generates a dense vector (embedding) for each text.
show_progress_bar=True → shows progress while encoding.
batch_size=32 → processes 32 texts at a time for efficiency.
...

print("✅ Embeddings shape:", len(embeddings))
```




```
c:\Users\abc\AppData\Local\Programs\Python\Python311\Lib\site-packages\tqdm\auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
Warning: You are sending unauthenticated requests to the HF Hub. Please set a H
F_TOKEN to enable higher rate limits and faster downloads.
Loading weights: 100%|██████████| 103/103 [00:00<00:00, 1826.51it/s]
BertModel LOAD REPORT from: sentence-transformers/all-MiniLM-L6-v2
Key | Status | |
-----+-----+---
embeddings.position_ids | UNEXPECTED | |

Notes:
- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
Batches: 100%|██████████| 219/219 [11:58<00:00, 3.28s/it]
✓ Embeddings shape: 7000
```

```
In [17]: import matplotlib.pyplot as plt
from sklearn.manifold import TSNE
import seaborn as sns

# Reduce embeddings to 2D
tsne = TSNE(n_components=2, random_state=42, perplexity=30)
X_tsne = tsne.fit_transform(embeddings)

# Scatter plot colored by topic
sns.scatterplot(
    x=X_tsne[:,0], y=X_tsne[:,1],
    hue=df["topic_label"], # use human-readable labels
    palette="tab10",
    alpha=0.7
)
plt.title("Semantic Embedding Clusters by Topic")
plt.show()
```



3. SEMANTIC SEARCH ENGINE

```
In [18]: from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

def semantic_search(query, top_k=5):
    query_emb = model.encode([query])
    #Converts the search query into an embedding vector using your SentenceTra

    scores = cosine_similarity(query_emb, embeddings)[0]
    #Computes cosine similarity between the query embedding and all paper embe

    top_idx = np.argsort(scores)[::-1][:top_k]
    #Sorts scores in descending order and picks the indices of the top-k most

    results = df.iloc[top_idx][["title", "topic_label"]].copy()
    #Retrieves the titles and topic labels of those top papers from your DataF

    results["score"] = scores[top_idx]
    #Adds similarity scores so you can see how close each match is.
```

```

    return results

# example
print(semantic_search("long context memory in LLMs", 5))

```

```

                                title \
6740 Developing Adaptive Context Compression Techni...
4529 Flux Attention: Context-Aware Hybrid Attention...
2161 MemoSight: Unifying Context Compression and Mu...
4010 KV Cache Offloading for Context-Intensive Tasks
4199 AsyncTLS: Efficient Generative LLM Inference w...

```

```

                                topic_label      score
6740 Multilingual & Speech  0.614558
4529 Multilingual & Speech  0.608411
2161 Reasoning & CoT  0.591219
4010 Multilingual & Speech  0.587048
4199 Multilingual & Speech  0.585296

```

```

In [19]: import matplotlib.pyplot as plt
         from sklearn.manifold import TSNE

         def visualize_search(query, top_k=5):
             # Run search
             results = semantic_search(query, top_k)

             # Get embeddings of top results + query
             query_emb = model.encode([query])
             top_embs = embeddings[results.index]

             combined = np.vstack([query_emb, top_embs])

             # Reduce to 2D
             tsne = TSNE(n_components=2, random_state=42, perplexity=5)
             X_tsne = tsne.fit_transform(combined)

             # Plot
             plt.figure(figsize=(8,6))
             plt.scatter(X_tsne[1:,0], X_tsne[1:,1], c="blue", label="Top Papers", alpha=0.5)
             plt.scatter(X_tsne[0,0], X_tsne[0,1], c="red", label="Query", marker="X", alpha=0.5)


             for i, title in enumerate(results["title"], start=1):
                 plt.text(X_tsne[i,0]+0.01, X_tsne[i,1]+0.01, title[:30]+"...", fontsize=8)

             plt.title(f"Semantic Search Visualization: '{query}'")
             plt.legend()
             plt.show()

             return results

         # Example usage
         visualize_search("long context memory in LLMs", 5)

```



Out[19]:

	title	topic_label	score
6740	Developing Adaptive Context Compression Techni...	Multilingual & Speech	0.614558
4529	Flux Attention: Context-Aware Hybrid Attention...	Multilingual & Speech	0.608411
2161	MemoSight: Unifying Context Compression and Mu...	Reasoning & CoT	0.591219
4010	KV Cache Offloading for Context-Intensive Tasks	Multilingual & Speech	0.587048
4199	AsyncTLS: Efficient Generative LLM Inference w...	Multilingual & Speech	0.585296

4. SIMILAR PAPER RETRIEVAL

```
In [20]: def find_similar_papers(paper_index, top_k=5):
#Takes the embedding vector of the paper at paper_index and reshapes it into
query_emb = embeddings[paper_index].reshape(1, -1)

#Computes similarity between this paper's embedding and all other embeddings
scores = cosine_similarity(query_emb, embeddings)[0]

...
Sorts scores in descending order.
[1:top_k+1] skips the first one (which would be the paper itself, similarity 1.0)
Picks the next top_k most similar papers.
```

```

'''
top_idx = np.argsort(scores)[::-1][1:top_k+1]

#Retrieves the titles and topic labels of those similar papers.
results = df.iloc[top_idx][["title", "topic_label"]].copy()
#Adds similarity scores to the results.
results["score"] = scores[top_idx]

return results

print(find_similar_papers(0, 5))

```

	title \	
2181	Adaptive Test-Time Compute Allocation for Reas...	
2865	Black-Box Optimization From Small Offline Data...	
5004	fastml: Guarded Resampling Workflows for Safer...	
810	Calibrating Scientific Foundation Models with ...	
6179	Test-Time Scaling Makes Overtraining Compute-0...	

	topic_label	score
2181	Multilingual & Speech	0.569100
2865	Neural Theory / Foundations	0.556702
5004	Neural Theory / Foundations	0.554789
810	Multilingual & Speech	0.549222
6179	Multilingual & Speech	0.540269

5. TREND SIGNAL DETECTION (VERY IMPORTANT)

In [21]: *# Convert time #Ensures the published column is treated as actual dates, not s*

```

df["published"] = pd.to_datetime(df["published"])

# Recent papers (last ~30 days assumption)
recent_cutoff = df["published"].max() - pd.Timedelta(days=30)
#Keeps only papers published within the last 30 days.
recent_df = df[df["published"] >= recent_cutoff]

print("Recent papers:", recent_df.shape)

```

Recent papers: (7000, 18)

In [22]: **import** matplotlib.pyplot **as** plt

```

# Count recent papers by topic
recent_topic_counts = recent_df["topic_label"].value_counts()


# Plot
plt.figure(figsize=(10,5))
recent_topic_counts.plot(kind="bar", color="skyblue")

plt.title("Recent Papers by Topic (Last 30 Days)", fontsize=14)

```

```
plt.xlabel("Topic", fontsize=12)
plt.ylabel("Number of Papers", fontsize=12)
plt.xticks(rotation=45, ha="right")
plt.grid(axis="y", linestyle="--", alpha=0.7)

plt.tight_layout()
plt.show()
```



6. EMERGING CLUSTER DETECTION

```
In [23]: from sklearn.cluster import KMeans

k = 8 # tunable

kmeans = KMeans(n_clusters=k, random_state=42)
clusters = kmeans.fit_predict(embeddings)

df["cluster"] = clusters
```

```
In [24]: import matplotlib.pyplot as plt


cluster_counts = df["cluster"].value_counts().sort_index()

plt.figure(figsize=(8,5))
cluster_counts.plot(kind="bar", color="skyblue")

plt.title("Distribution of Papers Across Clusters", fontsize=14)
plt.xlabel("Cluster ID", fontsize=12)
plt.ylabel("Number of Papers", fontsize=12)
plt.xticks(rotation=0)
plt.grid(axis="y", linestyle="--", alpha=0.7)

plt.tight_layout()
```


```
plt.show()
```



```
In [25]: from sklearn.manifold import TSNE
import seaborn as sns

tsne = TSNE(n_components=2, random_state=42, perplexity=30)
X_tsne = tsne.fit_transform(embeddings)

sns.scatterplot(
    x=X_tsne[:,0], y=X_tsne[:,1],
    hue=df["cluster"],
    palette="tab10",
    alpha=0.7
)
plt.title("t-SNE visualization of KMeans clusters")
plt.show()
```


```
In [26]: import matplotlib.pyplot as plt

cluster_counts = df["cluster"].value_counts().sort_index()

plt.figure(figsize=(8,5))
cluster_counts.plot(kind="bar", color="skyblue")

plt.title("Distribution of Papers Across Clusters", fontsize=14)
plt.xlabel("Cluster ID", fontsize=12)
plt.ylabel("Number of Papers", fontsize=12)
plt.xticks(rotation=0)
plt.grid(axis="y", linestyle="--", alpha=0.7)

plt.tight_layout()
plt.show()
```

8. SAVE ENRICHED DATASET

```
In [27]: output_path = r"C:\Users\abc\Downloads\arxiv_semantic_intelligence.csv"
df.to_csv(output_path, index=False)

print("✅ Saved:", output_path)
```

✅ Saved: C:\Users\abc\Downloads\arxiv_semantic_intelligence.csv

RESEARCH AGENT SYSTEM (FULL, END-TO-END)

🧠 Capabilities

Auto-extract core contributions

Identify method / problem / improvement

Compute novelty score

Compute impact score

Rank must-read papers

Generate research insights per cluster

```

In [1]: # =====
# PHASE 5 – RESEARCH AGENT SYSTEM
# =====

import pandas as pd
import numpy as np
import re
import ast
from collections import Counter
from sklearn.metrics.pairwise import cosine_similarity

# =====
# 1. LOAD DATA
# =====
file_path = r"C:\Users\abc\Downloads\arxiv_semantic_intelligence.csv"
df = pd.read_csv(file_path)

# =====
# 2. SAFE PARSE
# =====
def safe_eval(x):
    try:
        return ast.literal_eval(x) if isinstance(x, str) else []
    except:
        return []

df["tokens"] = df["tokens"].apply(safe_eval)

# =====
# 3. TEXT NORMALIZATION
# =====
df["text"] = df["clean_title"] + " " + df["clean_summary"]

# =====
# 4. CONTRIBUTION EXTRACTION (RULE-BASED NLP)
# =====

...
Uses regex patterns to detect phrases like “we introduce...” or “we propose...” to
...
def extract_contribution(text):
    text = text.lower()

    patterns = [
        r"we introduce (.+?)[\.\,]",
        r"we propose (.+?)[\.\,]",
        r"we present (.+?)[\.\,]",
        r"this paper introduces (.+?)[\.\,]",
        r"this work proposes (.+?)[\.\,]"
    ]

```

```

    for p in patterns:
        match = re.search(p, text)
        if match:
            return match.group(1)

    return "unknown"

df["contribution"] = df["text"].apply(extract_contribution)

# =====
# 5. METHOD / PROBLEM / IMPROVEMENT EXTRACTION
# =====

def extract_signal(text):
    text = text.lower()

    method = None
    problem = None
    improvement = None

    if "we introduce" in text or "we propose" in text:
        method = extract_contribution(text)

    if "problem" in text or "challenge" in text:
        problem = "identified problem"

    if "improve" in text or "outperform" in text:
        improvement = "performance improvement"

    return method, problem, improvement

df[["method", "problem", "improvement"]] = df["text"].apply(
    lambda x: pd.Series(extract_signal(x))
)

# =====
# 6. NOVELTY SCORE (SEMANTIC DISTANCE)
# =====

'''Loads embeddings (or regenerates them).

Computes the centroid of all embeddings.

Measures cosine similarity to the centroid → papers farther away are considered
novel.

Stores novelty as 1 - similarity.
'''
# Load embeddings from previous phase if available
try:
    import pickle
    with open("embeddings.pkl", "rb") as f:
        embeddings = pickle.load(f)
except:

```

```

from sentence_transformers import SentenceTransformer
model = SentenceTransformer("all-MiniLM-L6-v2")
embeddings = model.encode(df["text"].tolist(), show_progress_bar=True)

# Compute centroid
centroid = np.mean(embeddings, axis=0).reshape(1, -1)

# Distance from centroid = novelty
novelty_scores = cosine_similarity(embeddings, centroid).flatten()

df["novelty_score"] = 1 - novelty_scores

# =====
# 7. IMPACT SCORE
# =====

def compute_impact(row):
    score = 0

    text = row["text"]

    if "benchmark" in text:
        score += 2
    if "state of the art" in text or "sota" in text:
        score += 3
    if "outperform" in text:
        score += 2
    if "large scale" in text:
        score += 1
    if "open source" in text:
        score += 1

    return score

df["impact_score"] = df.apply(compute_impact, axis=1)

# =====
# 8. COMBINED RESEARCH SCORE
# =====

df["research_score"] = (
    0.6 * df["novelty_score"] +
    0.4 * (df["impact_score"] / df["impact_score"].max())
)

# =====
# 9. TOP PAPERS (GLOBAL)
# =====

top_papers = df.sort_values("research_score", ascending=False).head(20)

print("\n🔥 TOP RESEARCH PAPERS:")
print(top_papers[["title", "topic_label", "research_score"]].to_string(index=False))

```

```

# =====
# 10. CLUSTER INSIGHT GENERATION
# =====

def generate_cluster_insight(cluster_id):
    subset = df[df["cluster"] == cluster_id]

    words = []
    for tokens in subset["tokens"]:
        words.extend(tokens)

    top_words = [w for w, _ in Counter(words).most_common(10)]

    return ", ".join(top_words)

cluster_insights = {}

for c in sorted(df["cluster"].unique()):
    cluster_insights[c] = generate_cluster_insight(c)

print("\n🧠 CLUSTER INSIGHTS:")
for c, insight in cluster_insights.items():
    print(f"Cluster {c}: {insight}")

# =====
# 11. MUST-READ PER TOPIC
# =====

print("\n📖 MUST-READ PAPERS PER TOPIC:")

for topic in df["topic_label"].unique():
    subset = df[df["topic_label"] == topic]
    best = subset.sort_values("research_score", ascending=False).head(3)

    print(f"\n=== {topic} ===")
    print(best[["title", "research_score"]].to_string(index=False))

# =====
# 12. SAVE FINAL INTELLIGENCE DATASET
# =====

output_path = r"C:\Users\abc\Downloads\arxiv_research_agent.csv"
df.to_csv(output_path, index=False)

print("\n✅ FINAL DATASET SAVED:", output_path)

```

```

c:\Users\abc\AppData\Local\Programs\Python\Python311\Lib\site-packages\tqdm\auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
Warning: You are sending unauthenticated requests to the HF Hub. Please set a H
F_TOKEN to enable higher rate limits and faster downloads.
Loading weights: 100%|██████████| 103/103 [00:00<00:00, 7923.51it/s]
BertModel LOAD REPORT from: sentence-transformers/all-MiniLM-L6-v2
Key                | Status          | |
-----+-----+---+
embeddings.position_ids | UNEXPECTED     | |

Notes:
- UNEXPECTED          :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
Batches: 100%|██████████| 219/219 [10:24<00:00, 2.85s/it]

```

🔥 TOP RESEARCH PAPERS:

title	topic_label	research_score
LumaFlux: Lifting 8-Bit Worlds to HDR Reality with Physically-Guided Diffusion Transformers Knowledge & Continual Learning		0.755886
Corner Reflector Array Jamming Discrimination Using Multi-Dimensional Micro-Motion Features with Frequency Agile Radar	Neural Theory / Foundations	0.726627
Continuous Orthogonal Mode Decomposition: Haptic Signal Prediction in Tactile Internet	Neural Theory / Foundations	0.711159
Biomimetic causal learning for microstructure-forming phase transitions	Agents & Tool Use	0.708500
Multi-Agent Pathfinding with Non-Unit Integer Edge Costs via Enhanced Conflict-Based Search and Graph Discretization	Agents & Tool Use	0.704271
Do We Really Need to Approach the Entire Pareto Front in Many-Objective Bayesian Optimisation?	Agents & Tool Use	0.699136
Automatic dental superimposition of 3D intraorals and 2D photographs for human identification	Neural Theory / Foundations	0.695645
FlowPalm: Optical Flow Driven Non-Rigid Deformation for Geometrically Diverse Palmprint Generation	Neural Theory / Foundations	0.695204
Capturing Multivariate Dependencies of EV Charging Events: From Parametric Copulas to Neural Density Estimation	Neural Theory / Foundations	0.694936
Decoding High-Dimensional Finger Motion from EMG Using Riemannian Features and RNNs	Neural Theory / Foundations	0.685229
Greedy and Transformer-Based Multi-Port Selection for Slow Fluid Antenna Multiple Access	Agents & Tool Use	0.677901
Battery health prognosis using Physics-informed neural network with Quantum Feature mapping	Agents & Tool Use	0.674330
Robustness of Spatio-temporal Graph Neural Networks for Fault Location in Partially Observable Distribution Grids	Neural Theory / Foundations	0.672705
Channel-wise Retrieval for Multivariate Time Series Forecasting	AI Systems & Evaluation	0.672403
Parametric Interpolation of Dynamic Mode Decomposition for Predicting Nonlinear Systems	Agents & Tool Use	0.670993
$\text{F}^2\text{LP-AP}$: Fast & Flexible Label Propagation with Adaptive Propagation Kernel	Neural Theory / Foundations	0.670421
EEG-Based Emergency Braking Intensity Prediction Using Blind Source Separation	Neural Theory / Foundations	0.670299

CDSA-Net: Collaborative Decoupling of Vascular Structure and Background for High-Fidelity Coronary Digital Subtraction Angiography Neural Theory / Foundations
0.669803

Peer Review at Scale: The AAI-26 AI Review Pilot AI-Assisted Memory & Attention Systems
0.666087

ing Exact Covers Based on Decomposability Property A Parallel Approach to Counting Agents & Tool Use
0.664614

🧠 CLUSTER INSIGHTS:

Cluster 0: model, visual, image, task, language, reasoning, across, base, benchmark, vision

Cluster 1: model, policy, agent, task, learning, reward, base, learn, framework, training

Cluster 2: model, training, across, base, llm, memory, performance, time, achieve, large

Cluster 3: model, datum, base, dataset, use, method, feature, propose, framework, performance

Cluster 4: model, llm, human, agent, use, across, system, language, framework, large

Cluster 5: model, language, llm, reasoning, across, large, base, task, use, benchmark

Cluster 6: model, datum, method, base, framework, use, network, propose, learn, show

Cluster 7: agent, model, llm, system, task, framework, base, across, benchmark, tool

📖 MUST-READ PAPERS PER TOPIC:

=== Multilingual & Speech ===

	title	research_score
	Towards Real-Time ECG and EMG Modeling on μ NPUs	0.660676
	BioTrain: Sub-MB, Sub-50mW On-Device Fine-Tuning for Edge-AI on Biosignals	0.648990
	LoRaQ: Optimized Low Rank Approximation for 4-bit Quantization	0.628463

=== ML & Prediction Systems ===

title	research_score
LLM Safety From Within: Detecting Harmful Content with Internal Representations	0.646954
M-MiniGPT4: Multilingual VLLM Alignment via Translated Data	0.640878
DharmaOCR: Specialized Small Language Models for Structured OCR that outperform Open-Source and Commercial Baselines	0.630038

=== Multimodal / Vision ===

title	research_score
PEMANT: Persona-Enriched Multi-Agent Negotiation for Travel	

0.634689
Kernel-Smith: A Unified Recipe for Evolutionary Kernel Optimization
0.622898
SafeDream: Safety World Model for Proactive Early Jailbreak Detection
0.612616

=== Agents & Tool Use ===

title	research_score
Biomimetic causal learning for mic	
rostructure-forming phase transitions	0.708500
Multi-Agent Pathfinding with Non-Unit Integer Edge Costs via Enhanced Conflict-Based Search and Graph Discretization	0.704271
Do We Really Need to Approach the Entire Pareto Front in Many-Objective Bayesian Optimisation?	0.699136

=== Neural Theory / Foundations ===

title	research_score
Corner Reflector Array Jamming Discrimination Using Multi-Dimensional Micro-Motion Features with Frequency Agile Radar	0.726627
Continuous Orthogonal Mode Decomposition: Haptic Signal Prediction in Tactile Internet	0.711159
Automatic dental superimposition of 3D intraorals and 2D photographs for human identification	0.695645

=== AI Systems & Evaluation ===

title	research_score
Channel-wise Retrieval for Multivariate Time Series Forecasting	0.672403
Scalable and Adaptive Parallel Training of Graph Transformer on Large Graphs	0.651936
ORPHEAS: A Cross-Lingual Greek-English Embedding Model for Retrieval-Augmented Generation	0.642902

=== Reasoning & CoT ===

title	research_score
SURE: Synergistic Uncertainty-aware Reasoning for Multimodal Emotion Recognition in Conversations	0.614320
Decoupled Travel Planning with Behavior Forest	0.585614
CARO: Chain-of-Analogy Reasoning Optimization for Robust Content Moderation	0.577984

=== Reinforcement Learning ===

title	research_score
LiteResearcher: A Scalable Agentic RL Training Framework for Deep Research Agent	0.625535
A Pontryagin Method of Model-based Reinforcement Learning via Hamiltonian Actor-Critic	0.614513
Delayed Homomorphic Reinforcement Learning for Environments with Delayed Feedback	

eedback 0.614212

=== Memory & Attention Systems ===

	research_score	title
		AI-Assisted Peer Review at Scale: The AAAI-26 AI Review Pil
ot	0.666087	
		NFC based inventory control system for secure and efficient communicati
on	0.620306	
		"bot lane noob" Towards Deployment of NLP-based Toxicity Detectors in Video Gam
es	0.562754	


=== Knowledge & Continual Learning ===

	research_score	title
		LumaFlux: Lifting 8-Bit Worlds to HDR Reality with Physically-Guided D
iffusion Transformers	0.755886	
		Video-Robin: Autoregressive Diffusion Planning for Intent-Grounded Vide
o-to-Music Generation	0.659593	
		SAVE: A Generalizable Framework for Multi-Condition Single-Cell Generation with
Gene Block Attention	0.648142	

✓ FINAL DATASET SAVED: C:\Users\abc\Downloads\arxiv_research_agent.csv

```
In [2]: import matplotlib.pyplot as plt

plt.figure(figsize=(8,5))
df["research_score"].hist(bins=30, color="skyblue", edgecolor="black")
plt.title("Distribution of Research Scores", fontsize=14)
plt.xlabel("Research Score", fontsize=12)
plt.ylabel("Number of Papers", fontsize=12)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()
```



VISUALIZATION + LLM AGENT (ONE GO)

```
In [3]: # =====
# PHASE 6 – VISUALIZATION + LLM RESEARCH AGENT
# =====

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# =====
# 1. LOAD DATA
# =====
file_path = r"C:\Users\abc\Downloads\arxiv_research_agent.csv"
df = pd.read_csv(file_path)

# Clean datetime
df["published"] = pd.to_datetime(df["published"])

# Style (IMPORTANT)
sns.set(style="whitegrid", context="talk")
plt.rcParams["figure.figsize"] = (12, 6)

# =====
```

```

# 2. TOPIC DISTRIBUTION
# =====
plt.figure()
topic_counts = df["topic_label"].value_counts()

sns.barplot(
    x=topic_counts.values,
    y=topic_counts.index
)

plt.title("Research Topic Distribution")
plt.xlabel("Number of Papers")
plt.ylabel("Topic")
plt.tight_layout()
plt.show()

# =====
# 3. CLUSTER DISTRIBUTION
# =====
plt.figure()
sns.countplot(data=df, x="cluster")

plt.title("Cluster Distribution")
plt.xlabel("Cluster")
plt.ylabel("Count")
plt.tight_layout()
plt.show()

# =====
# 4. NOVELTY vs IMPACT (KEY GRAPH)
# =====
plt.figure()

sns.scatterplot(
    data=df,
    x="novelty_score",
    y="impact_score",
    hue="topic_label",
    alpha=0.7
)

plt.title("Research Landscape: Novelty vs Impact")
plt.xlabel("Novelty")
plt.ylabel("Impact")
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

# =====
# 5. TOP PAPERS (RESEARCH SCORE)
# =====
top = df.sort_values("research_score", ascending=False).head(15)

```

```

plt.figure()

sns.barplot(
    data=top,
    x="research_score",
    y="title"
)

plt.title("Top Research Papers")
plt.xlabel("Research Score")
plt.ylabel("Paper")
plt.tight_layout()
plt.show()

# =====
# 6. HEATMAP (TOPIC vs CLUSTER)
# =====
pivot = pd.crosstab(df["topic_label"], df["cluster"])

plt.figure(figsize=(14,8))
sns.heatmap(pivot, annot=True, fmt="d", cmap="coolwarm")

plt.title("Topic vs Cluster Heatmap")
plt.tight_layout()
plt.show()

# =====
# 7. TIME TREND
# =====
df["date"] = df["published"].dt.date

trend = df.groupby("date").size()

plt.figure()
sns.lineplot(x=trend.index, y=trend.values)

plt.title("Papers Over Time")
plt.xlabel("Date")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# =====
# 8. ADVANCED: FRONTIER PAPERS (HIGH NOVELTY + IMPACT)
# =====
frontier = df[
    (df["novelty_score"] > df["novelty_score"].quantile(0.8)) &
    (df["impact_score"] > df["impact_score"].quantile(0.8))
]

print("\n🚀 FRONTIER PAPERS:")
print(frontier[["title", "topic_label", "novelty_score", "impact_score"]].to_s

```

```

# =====
# 9. LLM-POWERED RESEARCH AGENT
# =====

from transformers import pipeline

# lightweight local model (no API needed)
summarizer = pipeline("text-generation", model="google/flan-t5-base")

def analyze_paper(text):
    prompt = f"""
    Analyze this research paper and extract:
    1. Main contribution
    2. Problem it solves
    3. Key method
    4. Why it matters

    Paper:
    {text[:1000]}
    """

    result = summarizer(prompt, max_length=200, do_sample=False)
    return result[0]["generated_text"]

# Apply to top papers only (efficiency)
top_papers = df.sort_values("research_score", ascending=False).head(5)

print("\n🧠 LLM RESEARCH INSIGHTS:\n")



for i, row in top_papers.iterrows():
    print(f"\n=== {row['title']} ===")
    insight = analyze_paper(row["text"])
    print(insight)


# =====
# 10. SAVE FINAL ENRICHED DATA
# =====

output_path = r"C:\Users\abc\Downloads\arxiv_final_visualized.csv"
df.to_csv(output_path, index=False)



print("\n✅ FINAL VISUAL + LLM DATASET SAVED:", output_path)


```





C:\Users\abc\AppData\Local\Temp\ipykernel_5668\314289281.py:88: UserWarning: Tight layout not applied. The left and right margins cannot be made large enough to accommodate all Axes decorations.
plt.tight_layout()





🚀 FRONTIER PAPERS:

title	topic_label	novelty_score	impact_score	
High-Dimensional Finger Motion from EMG Using Riemannian Features and RNNs	Decoding			N
neural Theory / Foundations	0.623530	7		
Untangling Data Purpose Compliance of Privacy Policies in Google Play Store	PrivSTRUCT: U			N
neural Theory / Foundations	0.643502	6		
Call-Chain-Aware LLM-Based Test Generation for Java Projects	ML & Prediction			
tion Systems	0.611351	5		
l-Constrained Neural Solver with Stochastic Physics-Informed Regularization	A Green-Integra			
Agents & Tool Use	0.618608	4		
Ideological Bias in LLMs' Economic Causal Reasoning	ML & Prediction Syst			
ems	0.659376	5		
image Super-Resolution at NTIRE 2026: Benchmark Results and Method Overview	The First Challenge on Remote Sensing Infrared I			
ledge & Continual Learning	0.696062	5		
able Evaluation Framework for GNN Fraud Ring Detection in Travel Networks	TRAVELFRAUDBENCH: A Configu			
AI Systems & Evaluation	0.620777	5		
for Accurate Prediction of Droplet Spreading Dynamics on Complex Surfaces	Droplet-LN0: Physics-Informed Laplace Neural Operators			
Agents & Tool Use	0.686362	5		
e detection of Controller Area Network attacks in the Internet of Vehicles	DAIRE: A lightweight AI model for real-tim			N
neural Theory / Foundations	0.616870	5		
t ₂ LP-AP: Fast & Flexible Label Propagation with Adaptive Propagation Kernel	F ^{textsuperscrip}			
Neural Theory / Foundations	0.598851	7		
rm Humans in Fraud Detection and Resistance to Motivated Investor Pressure	Large Language Models Outperfo			
Memory & Attention Systems	0.613486	4		
ral Networks for Fault Location in Partially Observable Distribution Grids	Robustness of Spatio-temporal Graph Neu			N
neural Theory / Foundations	0.602657	7		
Model for Adaptive Ultrasound-Guided Needle Insertion and Needle Tracking	A Vision-Language-Action			
neural Theory / Foundations	0.610672	6		
Graph-Theoretic Models for the Prediction of Molecular Measurements	Neural T			
heory / Foundations	0.619284	5		
sy Intermediate-Scale Quantum Computers: A Quantum Neural Network Approach	Option Pricing on Noi			
Agents & Tool Use	0.628581	5		
Based Emergency Braking Intensity Prediction Using Blind Source Separation	EEG-			N
neural Theory / Foundations	0.672720	6		
LoRaQ: Optimized Low Rank Approximation for 4-bit Quantization	Multili			
ngual & Speech	0.677068	5		

Towards Real-Time ECG and EMG Modeling on μ NPUs	Multilingual & Speech
h 0.730757 5	
Gaussian Splatting for Arbitrary-Scale Hyperspectral Image Super-Resolution Knowledge & Continual Learning	Voronoi-guided Bilateral 2D G
0.686266 5	
Farthest Point Sampling for Large-Scale Point Clouds via Pruning and Caching	FlashFPS: Efficient Fa
Multilingual & Speech 0.605285 4	
CDSA-Net: Collaborative Decoupling of Vascular Structure and Background for High-Fidelity Coronary Digital Subtraction Angiography	Neural Theory / Foundations
0.597820 7	
Forecast Sports Outcomes under Efficient Market Hypothesis: Theoretical and Experimental Analysis of Odds-Only and Generalised Linear Models	Neural Theory / Foundations
0.703961 4	
Flow Matching-Based Iterative Refinement for 3D Turbulent Flow Simulation	FlowRefiner:
Agents & Tool Use 0.659277 4	
Adaptive receptive field-based spatial-frequency feature reconstruction network for few-shot fine-grained image classification	Knowledge & Continual Learning
0.677674 5	
When Earth Foundation Models Meet Diffusion: An Application to Land Surface Temperature Super-Resolution Knowledge & Continual Learning	Neural Theory / Foundations
0.617809 4	
The CTLNet for Shanghai Composite Index Prediction	Neural Theory / Foundations
ns 0.642240 5	
Scalable and Adaptive Parallel Training of Graph Transformer on Large Graphs	Sc
AI Systems & Evaluation 0.642115 6	
SWNet: A Cross-Spectral Network for Camouflaged Weed Detection	Neural Theory / Foundations
0.734786 5	
Univariate Channel Fusion for Multivariate Time Series Classification	Neural Theory / Foundations
0.621378 5	
Corner Reflector Array Jamming Discrimination Using Multi-Dimensional Micro-Motion Features with Frequency Agile Radar	Neural Theory / Foundations
0.840674 5	
Stein Variational Black-Box Combinatorial Optimization	Agents & Tool Use
0.605090 6	
Lossless Compression via Chained Lightweight Neural Predictors with Information Inheritance	Agents & Tool Use
0.637406 5	
AI-Enabled Covert Channel Detection in RF Receiver Architectures	Neural Theory / Foundations
0.609877 6	
A Parallel Approach to Counting Exact Covers Based on Decomposability Property	A Pa
Agents & Tool Use 0.737319 5	
TOPCELL: Topology Optimization of Standard Cell via LLMs	Reinforcement Learning
0.622953 5	

			Magnitude
Is All You Need? Rethinking Phase in Quantum Encoding of Complex SAR Data			
Agents & Tool Use	0.636549	4	
			Bilevel Late Acc
eptance Hill Climbing for the Electric Capacitated Vehicle Routing Problem			
Agents & Tool Use	0.643251	6	
			Intelli
gent ROI-Based Vehicle Counting Framework for Automated Traffic Monitoring			N
neural Theory / Foundations	0.592247	4	
			Parametric In
terpolation of Dynamic Mode Decomposition for Predicting Nonlinear Systems			
Agents & Tool Use	0.747952	5	
Obtaining Partition Crossover masks using Statistical Linkage Learning for solving noised optimization problems with hidden variable dependency structure			
Agents & Tool Use	0.592416	5	
DynamicsLLM: a Dynamic Analysis-based Tool for Generating Intelligent Execution Traces Using LLMs to Detect Android Behavioural Code Smells			
Multimodal / Vision	0.610882	5	
Adaptive H-EFT-VA: A Provably Safe Trajectory Through the Trainability-Expressibility Landscape of Variational Quantum Algorithms			
Agents & Tool Use	0.611865	4	
			U
DAPose: Unsupervised Domain Adaptation for Low-Light Human Pose Estimation Knowledge & Continual Learning			
	0.598066	5	
			Battery health prognosis using Physics-informed neural network with Quantum Feature mapping
Agents & Tool Use	0.679440	6	
			PoreDi
T: A Scalable Generative Model for Large-Scale Digital Rock Reconstruction			N
neural Theory / Foundations	0.692688	4	
			Virtual Smart Metering in District Heating Networks via Heterogeneous Spatial-Temporal Graph Neural Networks
neural Theory / Foundations	0.639039	4	Ne
			FlowPalm: Optical Flow Driven Non-Rigid Deformation for Geometrically Diverse Palmprint Generation
neural Theory / Foundations	0.640155	7	N
			Continuous O
Orthogonal Mode Decomposition: Haptic Signal Prediction in Tactile Internet			N
neural Theory / Foundations	0.666747	7	
			Phys
Physics-guided surrogate learning enables zero-shot control of turbulent wings			
Reinforcement Learning	0.649155	5	
			Do We Really Need to
Approach the Entire Pareto Front in Many-Objective Bayesian Optimisation?			
Agents & Tool Use	0.646708	7	
			Persona-E ^{\$^2} : A Human-Grounded Dataset for Personality-Shaped Emotional Responses to Textual Events
ML & Prediction Systems	0.617192	4	
			A fast and Generic Energy-Shifting Transformer for Hybrid Monte Carlo Radiotherapy Calculation
neural Theory / Foundations	0.594216	4	N
			MATCHA: Efficient Deployment of Deep Neural Networks on Multi-Accelerator Heterogeneous Edge SoCs
Multilingual & Speech	0.624016	5	

The Traveling Thief Problem with Time Windows: Benchmarks and Heuristics	Neu
ral Theory / Foundations	0.678735 4
	CubeGra
ph: Efficient Retrieval-Augmented Generation for Spatial and Temporal Data	
AI Systems & Evaluation	0.636547 5
	Ti
me-Series Classification with Multivariate Statistical Dependence Features	
Agents & Tool Use	0.701040 5
DiffHDR: Re-Exposing LDR Videos with Video Diffusion Models Knowledge & Continual Learning	0.602762 5
	The
Character Error Vector: Decomposable errors for page-level OCR evaluation Knowledge & Continual Learning	0.620979 5
	NS-RGS: Newton-Sc
hulz based Riemannian gradient method for orthogonal group synchronization	
Agents & Tool Use	0.706025 4
	Automatic dental su
perimposition of 3D intraorals and 2D photographs for human identification	N
neural Theory / Foundations	0.789038 5
Thermal-GEMs: Generalized Models for Building Thermal Dynamics	Neural Theory
/ Foundations	0.602890 5
Channel-wise Retrieval for Multivariate Time Series Forecasting	AI Syste
ms & Evaluation	0.602153 7
	Multi-Agent Pathfinding with Non-Unit Inte
ger Edge Costs via Enhanced Conflict-Based Search and Graph Discretization	
Agents & Tool Use	0.655267 7
	SE-Enhanced ViT
and BiLSTM-Based Intrusion Detection for Secure IIoT and IoMT Environments	
Multilingual & Speech	0.650647 4
	Greedy and Tra
nsformer-Based Multi-Port Selection for Slow Fluid Antenna Multiple Access	
Agents & Tool Use	0.759465 5
HOIGS: Human-Object Interaction Gaussian Splatting Knowledge & Continual Learning	0.655790 5
	AC
ES: Who Tests the Tests? Leave-One-Out AUC Consistency for Code Generation	
ML & Prediction Systems	0.623978 5
	Understanding When Poisson Log-Norma
l Models Outperform Penalized Poisson Regression for Microbiome Count Data	N
neural Theory / Foundations	0.609553 4
	A Tsetlin M
achine-driven Intrusion Detection System for Next-Generation IoMT Security	N
neural Theory / Foundations	0.591001 5
	LumaFlux: Lifting
8-Bit Worlds to HDR Reality with Physically-Guided Diffusion Transformers Knowledge & Continual Learning	0.667217 8
Crystalite: A Lightweight Transformer for Efficient Crystal Modeling	
Agents & Tool Use	0.638885 5

Graph Neural Operator Towards Edge Deployability and Portability for Sparse-to-Dense, Real-Time Virtual Sensing on Irregular Grids			
Agents & Tool Use	0.631088	4	
Bridging Deep Learning and Integer Linear Programming: A Predictive-to-Prescriptive Framework for Supply Chain Analytics			N
Neural Theory / Foundations	0.604416	4	
			FSKD: Mon
ocular Forest Structure Inference via LiDAR-to-RGBI Knowledge Distillation			N
Neural Theory / Foundations	0.611269	5	
			Unsupervised 4D Flow MRI
Velocity Enhancement and Unwrapping Using Divergence-Free Neural Networks			
Agents & Tool Use	0.693500	5	
			Capturing Multivariate Dependencies of
f EV Charging Events: From Parametric Copulas to Neural Density Estimation			N
Neural Theory / Foundations	0.639708	7	
Biomimetic causal learning for microstructure-forming phase transitions			
Agents & Tool Use	0.662315	7	
			Differe
ntiable Initialization-Accelerated CPU-GPU Hybrid Combinatorial Scheduling			
Agents & Tool Use	0.597319	6	
			EpiScreen: Early Ep
ilepsy Detection from Electronic Health Records with Large Language Models			N
Neural Theory / Foundations	0.638357	5	
			Detecting low left ventricular ejection fracti
on from ECG using an interpretable and scalable predictor-driven framework			N
Neural Theory / Foundations	0.613288	4	
			KGroups: A Versatile Univariate Max-Relevance Min-R
edundancy Feature Selection Algorithm for High-dimensional Biological Data			N
Neural Theory / Foundations	0.637482	4	
			Optimi
zed Weighted Voting System for Brain Tumor Classification Using MRI Images			N
Neural Theory / Foundations	0.643430	5	
			Ne
iGAD: Augmenting Graph Anomaly Detection via Spectral Neighbor Information			N
Neural Theory / Foundations	0.659006	5	
FI-KAN: Fractal Interpolation Kolmogorov-Arnold Networks			Agents &
Tool Use	0.627917	4	
			Whitt
aker-Henderson smoother for long satellite image time series interpolation			
Agents & Tool Use	0.712673	5	
			MR-ImagenTime:
Multi-Resolution Time Series Generation through Dual Image Representations			N
Neural Theory / Foundations	0.650501	5	

Loading weights: 100%|██████████| 282/282 [00:00<00:00, 1640.73it/s]

The tied weights mapping and config for this model specifies to tie shared.weight to lm_head.weight, but both are present in the checkpoints, so we will NOT tie them. You should update the config with `tie\_word\_embeddings=False` to silence this warning

The model 'T5ForConditionalGeneration' is not supported for text-generation. Supported models are ['PeftModelForCausalLM', 'AfmoeForCausalLM', 'ApertusForCausalLM', 'ArceeForCausalLM', 'AriaTextForCausalLM', 'BambaForCausalLM', 'BartForCausalLM', 'BertLMHeadModel', 'BertGenerationDecoder', 'BigBirdForCausalLM', 'BigBirdPegasusForCausalLM', 'BioGptForCausalLM', 'BitNetForCausalLM', 'BlenderbotForCausalLM', 'BlenderbotSmallForCausalLM', 'BloomForCausalLM', 'BltForCausalLM', 'CamembertForCausalLM', 'LlamaForCausalLM', 'CodeGenForCausalLM', 'CohereForCausalLM', 'Cohere2ForCausalLM', 'CpmAntForCausalLM', 'CTRLLMHeadModel', 'CwmForCausalLM', 'Data2VecTextForCausalLM', 'DbrxForCausalLM', 'DeepseekV2ForCausalLM', 'DeepseekV3ForCausalLM', 'DiffLlamaForCausalLM', 'DogeForCausalLM', 'Dots1ForCausalLM', 'ElectraForCausalLM', 'Emu3ForCausalLM', 'ErnieForCausalLM', 'Ernie4_5ForCausalLM', 'Ernie4_5_MoeForCausalLM', 'Exaone4ForCausalLM', 'ExaoneMoeForCausalLM', 'FalconForCausalLM', 'FalconH1ForCausalLM', 'FalconMambaForCausalLM', 'FlexOlmoeForCausalLM', 'FuyuForCausalLM', 'GemmaForCausalLM', 'Gemma2ForCausalLM', 'Gemma3ForConditionalGeneration', 'Gemma3ForCausalLM', 'Gemma3nForConditionalGeneration', 'Gemma3nForCausalLM', 'GitForCausalLM', 'GlmForCausalLM', 'Glm4ForCausalLM', 'Glm4MoeForCausalLM', 'Glm4MoeLiteForCausalLM', 'GlmMoeDsaForCausalLM', 'GotOcr2ForConditionalGeneration', 'GPT2LMHeadModel', 'GPT2LMHeadModel', 'GPTBigCodeForCausalLM', 'GPTNeoForCausalLM', 'GPTNeoXForCausalLM', 'GPTNeoXJapaneseForCausalLM', 'GptOssForCausalLM', 'GPTJForCausalLM', 'GraniteForCausalLM', 'GraniteMoeForCausalLM', 'GraniteMoeHybridForCausalLM', 'GraniteMoeSharedForCausalLM', 'HeliumForCausalLM', 'HunYuanDenseV1ForCausalLM', 'HunYuanMoE V1ForCausalLM', 'Jais2ForCausalLM', 'JambaForCausalLM', 'JetMoeForCausalLM', 'Lfm2ForCausalLM', 'Lfm2MoeForCausalLM', 'LlamaForCausalLM', 'Llama4ForCausalLM', 'Llama4ForCausalLM', 'LongcatFlashForCausalLM', 'MambaForCausalLM', 'Mamba2ForCausalLM', 'MarianForCausalLM', 'MBartForCausalLM', 'MegatronBertForCausalLM', 'MiniMaxForCausalLM', 'MiniMaxM2ForCausalLM', 'MinistralForCausalLM', 'Ministral3ForCausalLM', 'MistralForCausalLM', 'MixtralForCausalLM', 'MllamaForCausalLM', 'ModernBertDecoderForCausalLM', 'MoshiForCausalLM', 'MptForCausalLM', 'MusicgenForCausalLM', 'MusicgenMelodyForCausalLM', 'MvpForCausalLM', 'NanoChatForCausalLM', 'NemotronForCausalLM', 'NemotronHForCausalLM', 'OlmoeForCausalLM', 'Olmoe2ForCausalLM', 'Olmoe3ForCausalLM', 'OlmoeHybridForCausalLM', 'OlmoeForCausalLM', 'OpenAIGPTLMHeadModel', 'OPTForCausalLM', 'PegasusForCausalLM', 'PersimmonForCausalLM', 'PhiForCausalLM', 'Phi3ForCausalLM', 'Phi4MultimodalForCausalLM', 'PhimoeForCausalLM', 'PLBartForCausalLM', 'ProphetNetForCausalLM', 'Qwen2ForCausalLM', 'Qwen2MoeForCausalLM', 'Qwen3ForCausalLM', 'Qwen3_5ForCausalLM', 'Qwen3_5MoeForCausalLM', 'Qwen3_5MoeForCausalLM', 'Qwen3_5ForCausalLM', 'Qwen3MoeForCausalLM', 'Qwen3NextForCausalLM', 'RecurrentGemmaForCausalLM', 'ReformerModelWithLMHead', 'RemBertForCausalLM', 'RobertaForCausalLM', 'RobertaPreLayerNormForCausalLM', 'RoCBertForCausalLM', 'RoFormerForCausalLM', 'RwkvForCausalLM', 'SeedOssForCausalLM', 'SmolLM3ForCausalLM', 'SolarOpenForCausalLM', 'StableLmForCausalLM', 'Starcoder2ForCausalLM', 'TrOCRForCausalLM', 'VaultGemmaForCausalLM', 'WhisperForCausalLM', 'XGLMForCausalLM', 'XLMWithLMHeadModel', 'XLMRobertaForCausalLM', 'XLMRobertaXLForCausalLM', 'XLNetLMHeadModel', 'xLSTMForCausalLM', 'XmodForCausalLM', 'YoutuForCausalLM', 'ZambaForCausalLM', 'Zamba2ForCausalLM'].

The following generation flags are not valid and may be ignored: ['temperature']. Set `TRANSFORMERS\_VERBOSITY=info` for more details.

Passing `generation\_config` together with generation-related arguments=({'do_sample', 'max_length'}) is deprecated and will be removed in future versions. Please

ase pass either a `generation\_config` object OR all generation parameters explicitly, but not both.

🧠 LLM RESEARCH INSIGHTS:

=== LumaFlux: Lifting 8-Bit Worlds to HDR Reality with Physically-Guided Diffusion Transformers ===

Both `max\_new\_tokens` (=256) and `max\_length` (=200) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Analyze this research paper and extract:

1. Main contribution
2. Problem it solves
3. Key method
4. Why it matters

Paper:

lumaflux lifting bit worlds to hdr reality with physically guided diffusion transformers the rapid adoption of hdr capable devices has created a pressing need to convert the bit standard dynamic range sdr content into perceptually and physically accurate bit high dynamic range hdr existing inverse tone mapping methods often rely on fixed tone mapping operators that struggle to generalize to real world degradations stylistic variations and camera pipelines frequently producing clipped highlights desaturated colors or unstable tone reproduction we introduce lumaflux a first physically and perceptually guided diffusion transformer dit for sdr to hdr reconstruction by adapting a large pretrained dit our lumaflux introduces a physically guided adaptation pga module that injects luminance spatial descriptors and frequency cues into attention through low rank residuals a perceptual cross modulation pcm layer that stabilizes chroma and texture via film conditioning from vision encoder feat

=== Corner Reflector Array Jamming Discrimination Using Multi-Dimensional Micro-Motion Features with Frequency Agile Radar ===

Both `max\_new\_tokens` (=256) and `max\_length` (=200) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Analyze this research paper and extract:

1. Main contribution
2. Problem it solves
3. Key method
4. Why it matters

Paper:

corner reflector array jamming discrimination using multi dimensional micro motion features with frequency agile radar this paper introduces a robust discrimination method for distinguishing real ship targets from corner reflector array jamming with frequency agile radar the key idea is to exploit the multidimensional micro motion signatures that separate rigid ships from non rigid decoys from range velocity maps we derive two new hand crafted descriptors mean weighted residual mwr and complementary contrast factor ccf and fuse them with deep features learned by a lightweight cnn an xgboost classifier then gives the final decision extensive simulations show that the hybrid feature set consistently outperforms state of the art alternatives confirming the superiority of the proposed approach

=== Continuous Orthogonal Mode Decomposition: Haptic Signal Prediction in Tactile Internet ===

Both `max\_new\_tokens` (=256) and `max\_length` (=200) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Analyze this research paper and extract:

1. Main contribution
2. Problem it solves
3. Key method
4. Why it matters

Paper:

continuous orthogonal mode decomposition haptic signal prediction in tactile internet the tactile internet demands sub millisecond latency and ultra high reliability as high latency or packet loss could lead to haptic control instability to address this we propose the mode domain architecture mda a bilateral predictive neural network architecture designed to restore missing signals on both the human and robot sides unlike conventional models that extract features implicitly from raw data mda utilizes a novel continuous orthogonal mode decomposition framework by integrating an orthogonality constraint we overcome the pervasive issue of mode overlapping found in state of the art decomposition methods experimental results demonstrate that this structured feature extraction achieves high prediction accuracies of human and robot furthermore the model achieves ultra low inference latency of ms significantly outperforming existing benchmarks and meeting the stringent real time requirements of

=== Biomimetic causal learning for microstructure-forming phase transitions ===

Both `max\_new\_tokens` (=256) and `max\_length` (=200) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Analyze this research paper and extract:

1. Main contribution
2. Problem it solves
3. Key method
4. Why it matters

Paper:

biomimetic causal learning for microstructure forming phase transitions non convex multi well energies in cell induced phase transitions give rise to fine scale microstructures low regularity transition layers and sharp interfaces all of which pose numerical challenges for physics informed learning to address this we propose biomimetic physics informed neural networks bio pinns for cell induced phase transitions in fibrous extracellular matrices the method converts the outward progression of cell mediated remodelling into a distance based training curriculum and couples it to uncertainty driven collocation that concentrates samples near evolving interfaces and tether forming regions the same uncertainty proxy provides a lower cost alternative to explicit second derivative regularization we also establish structural guarantees for the adaptive sampler including persistent coverage under gate expansion and quantitative near to far accumulation across single and multi cell benchmarks diver

=== Multi-Agent Pathfinding with Non-Unit Integer Edge Costs via Enhanced Conflict-Based Search and Graph Discretization ===

Both `max\_new\_tokens` (=256) and `max\_length` (=200) seem to have been set. `max\_new\_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/transformers/main/en/main_classes/text_generation)

Analyze this research paper and extract:

1. Main contribution
2. Problem it solves
3. Key method
4. Why it matters

Paper:

multi agent pathfinding with non unit integer edge costs via enhanced conflict based search and graph discretization multi agent pathfinding mapf plays a critical role in various domains traditional mapf methods typically assume unit edge costs and single timestep actions which limit their applicability to real world scenarios mapfr extends mapf to handle non unit costs with real valued edge costs and continuous time actions but its geometric collision model leads to an unbounded state space that compromises solver efficiency in this paper we propose mapfz a novel mapf variant on graphs with non unit integer costs that preserves a finite state space while offering improved realism over classical mapf to solve mapfz efficiently we develop cbs nic an enhanced conflict based search framework incorporating time interval based conflict detection and an improved safe interval path planning sipp algorithm additionally we propose bayesian optimization for graph design bogd a discretization met

✓ FINAL VISUAL + LLM DATASET SAVED: C:\Users\abc\Downloads\arxiv_final_visualized.csv

The "**Research Trends Analysis on AI and ML of 2026**" successfully demonstrates the power of an **automated intelligence pipeline** in navigating the vast and rapidly expanding landscape of academic literature. By processing 1,500 research papers through a multi-layered framework, this notebook provides a data-driven blueprint of current AI trajectories.

Key Findings in the 2026 AI Landscape

The analysis reveals that the research community is heavily focused on **applied autonomy and architectural efficiency**.

- **Dominant Research Domains:** The most active area of study is "**Agents & Tool Use**," accounting for 305 papers in the dataset, followed closely by "**Memory & Attention Systems**" with 236 papers. This indicates a significant shift toward creating AI systems that are not just passive models but active, agentic entities capable of tool manipulation and long-term reasoning.
- **High-Impact Frontiers:** The pipeline identified several "Frontier Papers" that represent the intersection of **high novelty and high impact**. Notable examples include research on **decoding high-dimensional finger motion from EMG** and **real-time ECG/EMG modeling on micro-neural processing units (μ NPUs)**, signaling a trend toward edge-deployed, privacy-preserving physiological AI.
- **The Rise of Specialized Intelligence:** Topic modeling highlights a robust diversification in AI sub-fields, with significant clusters emerging in **Multimodal/Vision** (141 papers) and **Knowledge & Continual Learning** (135 papers).

Methodological Success

The methodology employed in this notebook offers several critical advantages for modern researchers:

- **Semantic Over Keyword Search:** By utilizing **Sentence Transformer embeddings**, the system moves beyond simple keyword matching to "true meaning" retrieval, allowing for the discovery of related works that might use different terminology.
- **Automated Scoring Systems:** The introduction of **Novelty (semantic distance from the centroid)** and **Impact (signal detection for benchmarks and SOTA)** scores allows for the objective ranking of "must-read" papers, effectively filtering signal from noise.

- **LLM-Augmented Insights:** The integration of a **local LLM (Flan-T5)** to automatically extract contributions and problems solved by top papers demonstrates a scalable path for keeping up with the daily influx of new research without human fatigue.

Final Assessment

In conclusion, the 2026 AI research environment is characterized by an **evolution toward complex agentic behaviors and efficient hardware-aware architectures**. The tools developed in this notebook—from the **NMF-based topic distribution** to the **K-Means clustering** of semantic embeddings—provide the necessary infrastructure to transform raw ArXiv feeds into actionable research intelligence. This automated approach not only identifies current leaders in the field but also predicts emerging "frontier" ideas that will likely define the next wave of machine learning innovation.

In []: